

Vishesh Panghal

Department of Artificial Intelligence and Data Science
Poornima Institute of Engineering and Technology (PIET)
Jaipur, Rajasthan 302022, India
Email: pvt.panghal@gmail.com

September 24, 2026

To:

The Editor-in-Chief
Computer Methods and Programs in Biomedicine
Elsevier

Subject: Submission of Original Research Article

Dear Editor-in-Chief,

I am pleased to submit our original research manuscript entitled:

“Positive-Class Weighting Induces Systematic Probability Distortion in Deep Learning Models for Chest Radiography under Deidentified-Data Cohort Shift”

for publication consideration as an **Original Research Article** in *Computer Methods and Programs in Biomedicine*.

Significance and Alignment with CMPB Scope:

Clinical deep learning algorithms for chest radiography increasingly suffer from severe probability miscalibration when deployed over time. In biomedical image classification, positive-class weighting within binary cross-entropy (BCE) is standard practice to manage acute class imbalance (prevalence $< 5\%$). However, its algebraic interaction with model output logits has been widely conflated with intrinsic clinical feature drift.

Our study directly addresses the algorithmic and computational foundations of this failure:

- 1. Mathematical Derivation & Empirical Clustering:** We derive that positive loss weighting mathematically induces an additive constant logit displacement of $+\ln w$ in the population Bayes-optimal classifier. Across 16,604 evaluations spanning DenseNet-121 and ResNet-50 backbones on 40,523 radiographs from the BRAX benchmark, finite neural network logits cluster tightly around theoretical $\ln w$ (3.29 ± 0.81 vs. theoretical 3.07).
- 2. Structural Failure of Temperature Scaling:** We prove that standard temperature scaling (z/T) modifies logit variance without shifting the intercept, rendering it structurally incapable of correcting this intercept displacement ($\alpha \approx -2.96$).
- 3. Algorithmic $O(1)$ Analytic Recalibration:** We formulate a closed-form analytic logit offset ($z - \ln w$) that requires zero numerical optimization, zero validation tuning, and $O(1)$ constant

compute, eliminating 73.7% of stratum T_3 Brier error ($0.1079 \rightarrow 0.0284$, $\Delta\text{Brier} = -0.0795$ [95% CI: $-0.0893, -0.0697$]) and restoring calibration parity without model retraining.

4. **Budget-Ranked Selective Deferral:** We formalize the actual uncertainty-ranked selective triage pipeline, quantifying the clinical trade-off between calibrated risk estimation and recall-oriented screening referral.
5. **External Multi-Center Transportability Stress-Test:** We evaluate external validity on the official MIMIC-CXR-JPG frontal test split (3,403 images from 3,041 studies across 289 patients), establishing both the efficacy and the boundaries of zero-retraining recalibration under real-world domain shift.

Author Declarations and Compliance:

- This manuscript represents original work and has not been published previously, nor is it under simultaneous consideration for publication elsewhere.
- Publication is approved by the author and the host institution.
- No funding was received for this study, and the author declares no competing financial or non-financial interests.
- All analyzed benchmark cohorts (BRAX v1.1.0 and MIMIC-CXR-JPG v2.1.0) are fully anonymized, de-identified public datasets compliant with institutional ethics review and HIPAA standards.
- Complete code, checkpoints, and evaluation scripts supporting all findings are documented for deterministic replication.

Thank you for your time and consideration of our manuscript.

Sincerely,

Vishesh Panghal

Department of Artificial Intelligence and Data Science

Poornima Institute of Engineering and Technology (PIET), Jaipur, India

Corresponding Author: pvt.panghal@gmail.com