


Graphical Abstract

Positive-Class Weighting Induces Systematic Probability Distortion in Deep Learning Models for Chest Radiography under Deidentified-Date Cohort Shift

Vishesh Panghal



Highlights

Positive-Class Weighting Induces Systematic Probability Distortion in Deep Learning Models for Chest Radiography under Deidentified-Date Cohort Shift

Vishesh Panghal

- Positive loss weighting induces an additive population logit offset of $\log w$.
- Finite neural network logits empirically cluster tightly around theoretical $\log w$.
- Temperature scaling cannot correct this intercept-type calibration bias.
- Analytic offset ($z - \log w$) cuts 74% of Brier error without retraining.
- Selective triage exposes trade-offs between calibrated risk and screening referral.

Positive-Class Weighting Induces Systematic Probability Distortion in Deep Learning Models for Chest Radiography under Deidentified-Date Cohort Shift

Vishesh Panghal*

^aDepartment of Artificial Intelligence and Data Science, Poornima Institute of Engineering and Technology (PIET), Jaipur, 302022, Rajasthan, India

Abstract

Background and Objective: Clinical deep learning models for chest radiography frequently experience severe probability miscalibration under distribution shift. While positive-weighted binary cross-entropy (BCE) is standard practice to mitigate acute class imbalance, its mathematical distortion of output logit distributions has been conflated with feature drift. This study investigates the algebraic mechanism of loss-induced logit distortion, evaluates why standard post-hoc calibration fails, and formulates a zero-retraining recalibration and selective deferral pipeline.

Methods: We conducted a controlled 2×2 factorial experiment using two deep convolutional neural network architectures (DenseNet-121, ResNet-50) and two loss functions (unweighted vs. positive-weighted BCE) across three independent random seeds (12 models total) on 40,523 chest radiographs from BRAX. Models were trained on an anchor deployment stratum and evaluated across subsequent deidentified-date strata (T_1, T_2, T_3) with strict patient isolation. Recalibration strategies—temperature scaling, Platt scaling, and an $O(1)$ closed-form analytic logit offset ($z - \ln w$)—were evaluated alongside patient-clustered bootstrap ($B = 2000$). External validity was evaluated on MIMIC-CXR-JPG (3,403 radiographs, 289 patients).

Results: Positive-weighted BCE mathematically induces an additive population logit shift of $\ln w$, which temperature scaling structurally cannot resolve because scaling modifies logit dispersion without adjusting the intercept. In

*Corresponding author.

Email address: pvt.panghal@gmail.com (Vishesh Panghal)

empirical finite networks, logit differences clustered tightly around theoretical values (mean $\Delta z = 3.29 \pm 0.81$ vs. theoretical 3.07 for Pleural Effusion). On stratum T_3 , the zero-parameter analytic offset reduced Brier score by 73.7% ($\Delta\text{Brier} = -0.0795$ [95% CI: $-0.0893, -0.0697$], $0.1079 \rightarrow 0.0284$) and reduced intercept error from $\alpha = -2.96$ to -0.45 without validation tuning. On MIMIC-CXR-JPG, discrimination was pathology-dependent (AUROC 0.688 for Effusion), with analytic offset reducing Brier error for Cardiomegaly and Pneumonia. Uncertainty-based selective deferral reduced automated error rates to < 0.01 at 70% coverage.

Conclusions: Severe probability miscalibration in chest radiograph classification across held-out deidentified-date strata is largely an artifact of objective function weighting rather than intrinsic dataset drift. The closed-form analytic offset provides an immediate, constant-time $O(1)$ algorithmic recalibration that restores probability reliability across held-out deidentified-date cohorts without model refitting.

Keywords: Chest radiography, Probability calibration, Class imbalance, Distribution shift, Selective prediction, Deep learning, Clinical decision support

1. Introduction

Deep learning models for chest radiograph interpretation have achieved high discriminatory performance across common thoracic abnormalities [1, 2]. However, transferring models from stationary retrospective validation to longitudinal clinical deployment introduces substantial distribution shifts as patient populations, imaging protocols, and disease prevalence evolve over time [3, 4, 5].

While discriminatory ranking (AUROC) often exhibits gradual longitudinal attenuation, high-stakes medical decision support relies fundamentally on *calibrated posterior probabilities*: $P(Y = 1 \mid X = x)$ [6, 7]. Miscalibrated probabilities distort clinical decision cutoffs, degrade expected utility formulations, and compromise automated triage. Recent studies have reported rapid “temporal calibration decay,” asserting that frozen neural networks suffer progressive miscalibration over deployment periods [8, 9]. Concurrently, because thoracic datasets exhibit extreme class imbalance (prevalence often $< 5\%$), models are routinely trained using positive-weighted binary cross-entropy (BCE) [1, 10].

In this work, we demonstrate that these two phenomena are deeply entangled: the prevailing attribution of temporal calibration collapse to intrinsic feature drift is largely an artifact of uncorrected post-training calibration offsets under label distribution shifts [11, 12]. Under population risk minimization, training with positive loss weight $w > 1$ mathematically shifts the Bayes-optimal logit by an additive constant $+\ln w$ [13, 14]. Standard post-hoc calibration methods—most notably temperature scaling [15]—rescale logit variance (z/T) but cannot adjust the logit intercept. Consequently, weighted models continuously over-predict positive probabilities, producing severe Brier score and Expected Calibration Error (ECE) degradation that compounds as disease prevalence fluctuates over time.

1.1. Related Work

Probability Calibration and Class Imbalance: Deep neural networks frequently exhibit over-confidence due to unconstrained cross-entropy minimization [15]. While post-hoc recalibration techniques such as temperature scaling, histogram binning, and Platt scaling are standard [15, 16, 6], they were primarily formulated for balanced distributions. In clinical prediction modeling, Piccininni et al. [17] demonstrated that correcting class imbalance via resampling inherently distorts probability calibration, while plug-in corrections can recover valid absolute risks. In loss-based weighting, King and Zeng [13] established in classical logistic regression that choice-based sampling proportion adjustments induce an intercept shift of $\ln w$. Recently, Caplin, Martin, and Marx [14] analyzed loss-weighted deep learning through an economic lens, proving that class weighting is mathematically incompatible with raw posterior calibration and formulating an inverse machine learning framework to recover undistorted likelihoods on a static pneumonia task. However, their work was confined to single-split evaluations, leaving unexamined how loss-induced logit shifts behave under sequential temporal cohort shifts, whether standard temperature scaling fails, and how post-hoc recalibration performs under cross-institutional transfer.

Distribution Shift and Selective Prediction: Temporal dataset shift encompasses changes in feature distributions $P(X)$, conditional distributions $P(Y | X)$, and label priors $P(Y)$ [3, 12]. Alexandari et al. [11] established that bias-corrected calibration provides robust adaptation under label shift. Selective classification allows models to abstain on ambiguous inputs and route them to expert radiologists [18, 19, 20]. While selective triage studies focus primarily on abstention policies, the interaction between loss-induced

55 calibration distortion and downstream referral triage remains unexplored. In
 56 this work, we employ selective prediction as an operational lens to examine
 57 the real-world clinical trade-offs between calibrated probability estimation
 58 and screening triage.

59 1.2. Contributions

60 To resolve algorithmic loss bias from distribution shift across cohort strata,
 61 this paper provides:

- 62 1. **Controlled 2×2 Experimental Design:** We evaluate two clinical
 63 architectures (DenseNet-121, ResNet-50) and two loss objectives (stan-
 64 dard vs. positive-weighted BCE) across three independent seeds (12
 65 models total) on 40,523 radiographs from BRAX [2], tracking discrim-
 66 ination and calibration across three held-out deidentified-date strata
 67 (T_1, T_2, T_3) with strict patient isolation.
- 68 2. **Mathematical Derivation and Empirical Clustering of Logit**
 69 **Shift:** We prove that positive-weighted BCE introduces a constant
 70 population logit offset $\ln w$ (Section S1 of the Supplementary Material).
 71 Across 16,604 evaluations, finite neural network logits cluster tightly
 72 around theoretical $\ln w$ (mean $\Delta z = 3.29 \pm 0.81$ vs. theoretical 3.07
 73 for Pleural Effusion). Crucially, we prove that temperature scaling
 74 structurally fails to correct this intercept displacement ($\alpha \approx -2.96$),
 75 whereas a closed-form analytic offset ($z - \ln w$) eliminates 73.7% of
 76 stratum T_3 Brier error ($0.1079 \rightarrow 0.0284$) and restores calibration parity
 77 without model retraining.
- 78 3. **Leak-Free Cross-Fitting and Patient-Clustered Bootstrap:** We
 79 implement a 5-fold patient-clustered cross-fitting protocol on T_1 and
 80 evaluate out-of-sample generalization on T_2 and T_3 backed by 2,000-
 81 replicate patient-clustered bootstrap intervals.
- 82 4. **Operational Selective Prediction and Cross-Institutional Trans-**
 83 **portability:** We formalize budget-constrained uncertainty-ranked selec-
 84 tive deferral, quantifying the trade-off between calibrated risk estimation
 85 and screening triage. We stress-test external transportability on MIMIC-
 86 CXR-JPG (3,403 images, 289 patients), delineating the capabilities and
 87 boundaries of zero-retraining recalibration under real-world domain
 88 shift.

Table 1: Cohort Demographics, Radiographic Projections, and Target Condition Prevalence across Stratified Cohorts

Cohort Stratum	Images (N)	Patients (N)	Studies (N)	Sex (% M / F)	Age (Mean $\pm$ SD)	Views (% PA / AP / Lat / Unk)	Pleural Effusion N (%)	Cardiomegaly N (%)	Pneumonia N (%)	Edema N (%)
Train (Pre-deployment)	25,123	11,265	14,315	49.3% / 50.7%	46.0 $\pm$ 24.3	34.0% / 13.5% / 0.0% / 27.2%	1,116 (4.44%)	2,398 (9.55%)	456 (1.82%)	23 (0.09%)
Val (Pre-deployment)	7,098	3,281	4,012	49.8% / 50.2%	46.2 $\pm$ 24.1	34.7% / 12.2% / 0.0% / 27.7%	295 (4.16%)	681 (9.59%)	128 (1.80%)	8 (0.11%)
Test T1 (2015)	1,879	913	1,039	47.8% / 52.2%	44.5 $\pm$ 23.9	34.4% / 10.2% / 0.0% / 29.0%	99 (5.27%)	147 (7.82%)	50 (2.66%)	4 (0.21%)
Test T2 (2016)	3,987	1,832	2,249	48.5% / 51.5%	46.1 $\pm$ 23.8	34.6% / 12.0% / 0.0% / 27.6%	177 (4.44%)	422 (10.58%)	79 (1.98%)	13 (0.33%)
Test T3 (2017)	2,436	1,164	1,382	50.9% / 49.1%	46.0 $\pm$ 24.6	34.4% / 12.4% / 0.0% / 28.1%	85 (3.49%)	258 (10.59%)	50 (2.05%)	0 (0.00%)
Test All (Pooled)	8,302	3,807	4,670	49.0% / 51.0%	45.7 $\pm$ 24.1	34.5% / 11.7% / 0.0% / 28.1%	361 (4.35%)	827 (9.96%)	179 (2.16%)	17 (0.20%)

2. Materials and Methods

2.1. Dataset and Cohort Stratification

We utilize the Brazilian Chest X-ray (BRAX) v1.1.0 benchmark [2]. BRAX contains 40,967 radiographs labeled with 14 thoracic findings extracted from Portuguese reports via a validated NLP pipeline. After excluding 444 radiographs from 89 patients whose examinations crossed partition boundaries, the analytic cohort comprised 40,523 radiographs acquired at Hospital Israelita Albert Einstein in São Paulo, Brazil.

In compliance with HIPAA de-identification standards, timestamps are protected as fictitious 8-digit calendar integers (YYYYMMDD) preserving intra-patient acquisition intervals. Because cross-patient calendar synchronization is unverified, partitions represent sequential deidentified-date strata rather than synchronized calendar years. Studies are partitioned chronologically around a strict model freeze boundary ($T_{\text{freeze}} = 2015-07-01$): training studies before 2013-09-01 ($N = 25,123$ images, 11,265 patients, 14,315 studies); validation studies between 2013-09-01 and 2015-06-30 ($N = 7,098$ images, 3,281 patients, 4,012 studies); and held-out test studies between 2015-07-01 and 2017-12-17 ($N = 8,302$ images, 3,807 patients, 4,670 studies) partitioned into three consecutive strata: T_1 ($N = 1,879$), T_2 ($N = 3,987$), and T_3 ($N = 2,436$). Strict patient isolation is enforced ($\text{Patients}_{\text{train}} \cap \text{Patients}_{\text{val}} \cap \text{Patients}_{\text{test}} = \emptyset$). Complete cohort characteristics are summarized in Table 1.

2.2. External Stress-Testing Cohort (MIMIC-CXR-JPG)

To assess whether the mathematical calibration mechanism survives cross-institutional domain shift, frozen BRAX models were externally evaluated on all eligible frontal radiographs (AP and PA views) from the official patient-partitioned MIMIC-CXR-JPG test split [21, 22]. The cohort comprises $N = 3,403$ frontal radiographs ($N = 3,041$ diagnostic studies across $N = 289$ unique patients; AP: 2,405 [70.67%], PA: 998 [29.33%]). Ground-truth labels were extracted using the official CheXpert labeler under the uncertainty-zero (U -zero) policy: Pleural Effusion 32.18% ($N = 1,095$ positive radiographs),

119 Cardiomegaly 26.33% ($N = 896$), Edema 21.33% ($N = 726$), and Pneumonia
 120 10.05% ($N = 342$). Patient-clustered non-parametric bootstrap inference
 121 ($B = 1,000$ replicates) accounts for repeated within-patient examinations. No
 122 model fine-tuning or retraining was performed.

123 2.3. Architectures and Training Objectives

124 We evaluate two standard clinical vision architectures initialized from
 125 ImageNet pre-trained weights:

- 126 • **DenseNet-121** [23]: Employs dense feature concatenation, standard
 127 for thoracic pathology benchmarks [1].
- 128 • **ResNet-50** [24]: Employs residual skip connections as an independent
 129 structural comparator.

130 Each architecture is trained under two distinct loss objectives:

131 1. Unweighted Binary Cross-Entropy (Standard BCE):

$$\mathcal{L}_{\text{std}}(y, \hat{p}) = -[y \log \hat{p} + (1 - y) \log(1 - \hat{p})] \quad (1)$$

132 where $\hat{p} = \sigma(z) = 1/(1 + e^{-z})$ is the sigmoid probability from raw logit
 133 z .

134 2. Positive-Weighted Binary Cross-Entropy (Weighted BCE):

$$\mathcal{L}_{\text{pos}}(y, \hat{p}) = -[w \cdot y \log \hat{p} + (1 - y) \log(1 - \hat{p})] \quad (2)$$

135 where $w = (N_{\text{total}} - N_{\text{pos}})/N_{\text{pos}}$ represents the inverse positive prevalence
 136 in the training set (e.g., $w = 21.51$ for Pleural Effusion, $w = 9.48$ for
 137 Cardiomegaly, $w = 54.09$ for Pneumonia).

138 Models are trained with the AdamW optimizer (learning rate 10^{-4} , weight
 139 decay 10^{-2}) and a Cosine Annealing schedule over 15 epochs (batch size
 140 64, AMP fp16, 512×512 resolution). Pleural Effusion was pre-specified
 141 as the primary target pathology due to its stable evaluability across all
 142 longitudinal held-out strata ($T_1 = 1,879, T_2 = 3,987, T_3 = 2,436$); Car-
 143 diomegaly, Pneumonia, and Edema served as secondary mechanistic vali-
 144 dation targets. Checkpoints were selected by highest validation AUROC.
 145 Three random seeds ($S \in \{42, 123, 2026\}$) were trained per configuration (12
 146 models total). For unified benchmarking, a 3-seed probability ensemble was
 147 constructed by averaging predicted probabilities: $\bar{p}_{\text{raw}}(x) = \frac{1}{3} \sum_{s \in S} \sigma(z_s(x))$

148 and $\bar{p}_{\text{corrected}}(x) = \frac{1}{3} \sum_{s \in S} \sigma(z_s(x) - \ln w)$. The empirical logit offset was
 149 evaluated across all 16,604 ensemble-paired evaluations (8,302 images $\times$ 2
 150 architectures; 49,812 seed-level pairs):

$$\Delta z(x) = \bar{z}_w(x) - \bar{z}_u(x) = \frac{1}{3} \sum_{s \in S} (z_{w,s}(x) - z_{u,s}(x)) \quad (3)$$

151 2.4. Mathematical Mechanics of Loss Weighting and Calibration

152 Under positive-weighted binary cross-entropy, the conditional expected
 153 risk with positive weight $w > 0$ is:

$$\mathbb{E}_{Y|X}[\mathcal{L}_{\text{pos}}(Y, \hat{p})] = -w \cdot P(Y = 1 | X) \ln \hat{p} - P(Y = 0 | X) \ln(1 - \hat{p}) \quad (4)$$

154 Setting the first-order condition $\frac{\partial}{\partial z} \mathbb{E}[\mathcal{L}_{\text{pos}}] = 0$ yields the optimal predicted
 155 probability:

$$\hat{p}^*(x) = \frac{w \cdot P(Y = 1 | X = x)}{w \cdot P(Y = 1 | X = x) + P(Y = 0 | X = x)} \quad (5)$$

156 Converting to log-odds yields the Bayes-optimal logit:

$$z^*(x) = \ln \left(\frac{P(Y = 1 | X = x)}{P(Y = 0 | X = x)} \right) + \ln w = z_{\text{true}}(x) + \ln w \quad (6)$$

157 (see Section S1 of the Supplementary Material for the complete algebraic
 158 proof). Equation (6) derives that positive loss weighting shifts the model’s
 159 population Bayes-optimal logit by $+\ln w$, connecting classical choice-based
 160 sampling corrections [13] with the class-weight likelihood recovery formulation
 161 of Caplin et al. [14]. In finite neural networks, empirical logits fluctuate
 162 around this theoretical attractor.

163 2.4.1. Limitation of Temperature Scaling

164 Temperature scaling [15] adjusts probabilities via a single scalar T : $p_{\text{temp}} =$
 165 $\sigma(z/T)$. In the Cox calibration framework [6, 7]:

$$\text{logit}(P(Y = 1 | X)) = \alpha + \beta \cdot \text{logit}(\hat{p}) \quad (7)$$

166 where $\alpha = 0$ represents perfect calibration-in-the-large and $\beta = 1$ denotes ideal
 167 slope. Because positive weighting induces an additive shift $z \approx z_{\text{true}} + \ln w$,
 168 the uncorrected logit suffers an acute intercept discrepancy ($\alpha \approx -\ln w \ll 0$).
 169 Temperature scaling multiplies logits by $1/T$, adjusting slope $\beta \rightarrow \beta \cdot T$,
 170 but is mathematically incapable of adjusting the additive constant α . In
 171 low-prevalence clinical domains where negative cases dominate, the optimizer
 172 is constrained by conflicting slope and intercept pressures, typically selecting
 173 $T^* \approx 1.0$ without resolving the underlying intercept displacement.

174 2.4.2. Post-Hoc Recalibration Suite and Computational Complexity

175 We evaluate four post-hoc recalibration functions:

- 176 1. **Temperature Scaling (T^*):** $p = \sigma(z/T^*)$, optimized via bounded
177 Brent minimization on $[0.01, 50.0]$ with an explicit monotonicity con-
178 straint; incurs $O(N_{\text{val}})$ optimization and $O(1)$ inference.
- 179 2. **Analytic Offset ($z - \ln w$):** Directly subtracts the theoretical logit
180 shift: $p = \sigma(z - \ln w)$. Requires zero validation data and zero numer-
181 ical optimization ($O(1)$ training and inference), resetting $\alpha \approx 0$ by
182 construction while preserving model discrimination and slope.
- 183 3. **Fitted Offset ($z + a^*$):** Optimizes a scalar intercept $a^* \in [-10.0, 10.0]$
184 on validation log-likelihood, holding slope fixed at 1.0 ($O(N_{\text{val}})$ fitting,
185 $O(1)$ inference).
- 186 4. **Platt Scaling ($a + bz$):** Logistic regression on logits with a non-negative
187 slope constraint ($b \geq 0$) ($O(k \cdot N_{\text{val}})$ fitting, $O(1)$ inference).

188 2.5. Cross-Fitting and Resampling Evaluation Protocol

189 On deployment stratum T_1 (2015-H2), calibrator parameters are estimated
190 using **5-fold patient-clustered cross-fitting** (all studies from any given
191 patient assigned to a single fold). Calibrators fit on T_1 are frozen and
192 applied strictly out-of-sample on later strata T_2 (2016) and T_3 (2017). All
193 95% confidence intervals are computed using $B = 2,000$ patient-clustered
194 bootstrap resamples with replacement ($B = 1,000$ for external MIMIC
195 validation). Paired calibrator contrasts are evaluated on identical bootstrap
196 resamples.

197 2.6. Selective Prediction and Deferral Policy

198 In operational deployment, models can selectively defer ambiguous ex-
199 aminations to expert radiologists [18, 19, 20]. Rather than developing novel
200 rejection algorithms, our selective prediction formulation evaluates how loss
201 weighting and calibration impact clinical referral efficiency and automated
202 error rates.

203 Operating decision cutoffs t^* are determined strictly on development
204 stratum T_1 using Youden’s J index ($t^* \approx 0.014$ for unweighted, $t^* \approx 0.022$ for
205 analytic offset, $t^* \approx 0.33$ for raw weighted models). We define a threshold-
206 relative uncertainty metric based on piecewise distance from t^* :

$$u(x) = 1 - \frac{|\hat{p}(x) - t^*|}{\mathbb{I}(\hat{p}(x) \geq t^*)(1 - t^*) + \mathbb{I}(\hat{p}(x) < t^*)t^*} \in [0, 1] \quad (8)$$




Figure 1: Multi-calibrator evaluation on Pleural Effusion for the DenseNet-121 3-seed ensemble across sequential deidentified-date strata (T_1, T_2, T_3), directly visualizing Table 3. (a) Brier score and (b) Log-loss. Temperature scaling fails to reduce probability error in weighted models, whereas Analytic Offset ($z - \log w$) and Platt Scaling reduce Brier score by 73.7% ($\Delta\text{Brier} = -0.0795$), matching unweighted models.

where $u(x) \rightarrow 1$ indicates maximal ambiguity ($\hat{p}(x) \approx t^*$) and $u(x) \rightarrow 0$ indicates decisive confidence. At coverage budget $C \in [0.5, 1.0]$, the system retains the fraction C of cases with lowest uncertainty ($N_{\text{auto}} = \lceil C \cdot N \rceil$) and refers the remaining $1 - C$ to radiologist review ($N_{\text{ref}} = N - N_{\text{auto}}$).

We track triage outcomes across five distinct metrics: (i) retained positives ($N_{\text{pos, ret}}$); (ii) automated true positives ($TP_{\text{auto}} = \sum_{i \in \text{Auto}} \mathbb{I}(y_i = 1 \wedge \hat{y}_i = 1)$); (iii) automated false negatives ($FN_{\text{auto}} = \sum_{i \in \text{Auto}} \mathbb{I}(y_i = 1 \wedge \hat{y}_i = 0)$, where $TP_{\text{auto}} + FN_{\text{auto}} = N_{\text{pos, ret}}$); (iv) retained sensitivity ($\text{Sens}_{\text{ret}} = TP_{\text{auto}} / N_{\text{pos, ret}}$); and (v) full-cohort triage sensitivity ($\text{Sens}_{\text{triage}} = 1 - FN_{\text{auto}} / N_{\text{pos, total}}$), assuming deferred cases are correctly diagnosed by expert radiologists.

3. Experimental Results

3.1. Experimental Matrix: Discrimination vs. Calibration Trajectories

Table 2 presents the comprehensive evaluation of the 2×2 experimental design comprising two architectures and two loss objectives (with three independently trained random seed replications per configuration, 12 models total) for Pleural Effusion, Cardiomegaly, and Pneumonia across chronological strata T_1, T_2, T_3 .

3.1.1. Discrimination Attenuation vs. Calibration Trajectories

Both architectures exhibit moderate discriminatory attenuation across subsequent deidentified-date cohort strata. For DenseNet-121 on Pleural

Algorithm 1 Budget-Constrained Uncertainty-Ranked Selective Deferral Pipeline

Require: Cohort radiographs $\{x_i\}_{i=1}^N$, trained neural network f_θ , positive loss weight w , clinical operating threshold t^* , target coverage budget $C \in (0, 1]$.

Ensure: Automated decisions $\{\hat{y}_i : i \in \mathcal{S}\}$ and referred examination set $\mathcal{D}$.

```

1: Phase 1: Analytic Recalibration and Uncertainty Scoring
2: for  $i = 1$  to  $N$  do
3:   Compute raw network logit:  $z_i \leftarrow f_\theta(x_i)$ 
4:   if model was trained with positive class weight  $w > 1$  then
5:     Apply constant-time closed-form offset:  $\tilde{z}_i \leftarrow z_i - \ln(w)$ 
6:   else
7:      $\tilde{z}_i \leftarrow z_i$ 
8:   end if
9:   Compute calibrated probability:  $\hat{p}_i \leftarrow \sigma(\tilde{z}_i) = \frac{1}{1 + \exp(-\tilde{z}_i)}$ 
10:  Compute threshold-relative normalizer:
11:     $s_i \leftarrow \begin{cases} 1 - t^*, & \text{if } \hat{p}_i \geq t^* \\ t^*, & \text{otherwise} \end{cases}$ 
12:  Compute normalized threshold-proximity uncertainty:
13:     $u_i \leftarrow 1 - \min\left(1, \frac{|\hat{p}_i - t^*|}{s_i}\right) \in [0, 1]$ 
14: end for
15: Phase 2: Budget-Constrained Triage Partitioning
16: Compute automation quota:  $k \leftarrow \min(N, \max(1, \lceil C \cdot N \rceil))$ 
17: Sort cohort indices by ascending uncertainty:  $\pi \leftarrow \text{argsort}(\{u_1, \dots, u_N\})$ 
18: Partition cohort into automated and deferred sets:
19:    $\mathcal{S} \leftarrow \{\pi(1), \pi(2), \dots, \pi(k)\}$    {Retained automated cohort}
20:    $\mathcal{D} \leftarrow \{\pi(k+1), \dots, \pi(N)\}$    {Referred to expert radiologist review}
21: Phase 3: Decision Assignment
22: for  $i \in \mathcal{S}$  do
23:   Assign automated binary classification:  $\hat{y}_i \leftarrow \mathbb{I}(\hat{p}_i \geq t^*)$ 
24: end for
25: return Automated decisions  $\{\hat{y}_i\}_{i \in \mathcal{S}}$ , deferred cases  $\mathcal{D}$ .
```

228 Effusion, AUROC decreases from 0.9335 ± 0.0052 at T_1 to 0.8731 ± 0.0120 at T_3
229 ($\Delta\text{AUROC} = -0.0604$); ResNet-50 exhibits an identical downward trajectory
230 ($0.9295 \pm 0.0113 \rightarrow 0.8608 \pm 0.0190$, $\Delta\text{AUROC} = -0.0687$; Table 2). This

Table 2: Experimental Matrix Evaluation: Primary Longitudinal Performance across Architectures, Loss Objectives, and Thoracic Pathologies (Strata $T_1 \rightarrow T_3$). Values represent mean $\pm$ standard deviation across three independently trained random seeds ($S \in \{42, 123, 2026\}$). Complete 17-column trajectories across all intermediate deployment strata (T_1, T_2, T_3) and log-loss metrics are detailed in Supplementary Table S3.

Architecture	Objective	Pathology	AUROC ($T_1 \rightarrow T_3$)	Δ AUROC	Brier ($T_1 \rightarrow T_3$)	Δ Brier	ECE T_3
DenseNet-121	Unweighted BCE	Pleural Effusion	0.934 $\rightarrow$ 0.873	-0.060	0.037 $\rightarrow$ 0.028	-0.009	0.0090
DenseNet-121	Unweighted BCE	Cardiomegaly	0.908 $\rightarrow$ 0.902	-0.006	0.053 $\rightarrow$ 0.068	+0.015	0.0210
DenseNet-121	Unweighted BCE	Pneumonia	0.847 $\rightarrow$ 0.865	+0.017	0.024 $\rightarrow$ 0.019	-0.005	0.0095
DenseNet-121	Unweighted BCE	Edema*	0.540 $\rightarrow$ -	-	0.002 $\rightarrow$ 0.000	-0.002	0.0029
DenseNet-121	Pos-Weighted BCE	Pleural Effusion	0.906 $\rightarrow$ 0.839	-0.067	0.108 $\rightarrow$ 0.125	+0.017	0.2227
DenseNet-121	Pos-Weighted BCE	Cardiomegaly	0.874 $\rightarrow$ 0.858	-0.016	0.126 $\rightarrow$ 0.136	+0.010	0.1969
DenseNet-121	Pos-Weighted BCE	Pneumonia	0.818 $\rightarrow$ 0.842	+0.024	0.131 $\rightarrow$ 0.130	-0.002	0.2425
DenseNet-121	Pos-Weighted BCE	Edema*	0.612 $\rightarrow$ -	-	0.061 $\rightarrow$ 0.064	+0.002	0.1274
ResNet-50	Unweighted BCE	Pleural Effusion	0.930 $\rightarrow$ 0.861	-0.069	0.037 $\rightarrow$ 0.028	-0.009	0.0132
ResNet-50	Unweighted BCE	Cardiomegaly	0.890 $\rightarrow$ 0.879	-0.011	0.057 $\rightarrow$ 0.071	+0.014	0.0175
ResNet-50	Unweighted BCE	Pneumonia	0.843 $\rightarrow$ 0.880	+0.037	0.024 $\rightarrow$ 0.019	-0.005	0.0106
ResNet-50	Unweighted BCE	Edema*	0.556 $\rightarrow$ -	-	0.002 $\rightarrow$ 0.000	-0.002	0.0017
ResNet-50	Pos-Weighted BCE	Pleural Effusion	0.879 $\rightarrow$ 0.819	-0.060	0.126 $\rightarrow$ 0.139	+0.014	0.2751
ResNet-50	Pos-Weighted BCE	Cardiomegaly	0.830 $\rightarrow$ 0.799	-0.031	0.142 $\rightarrow$ 0.152	+0.011	0.2338
ResNet-50	Pos-Weighted BCE	Pneumonia	0.835 $\rightarrow$ 0.823	-0.012	0.167 $\rightarrow$ 0.169	+0.002	0.3253
ResNet-50	Pos-Weighted BCE	Edema*	0.549 $\rightarrow$ -	-	0.069 $\rightarrow$ 0.075	+0.005	0.1670

confirms that discriminatory attenuation occurs gradually across subsequent cohort strata regardless of loss objective.

In stark contrast, probability calibration diverges sharply by loss objective. For unweighted models, Brier scores are low and actually improve as prevalence falls ($0.0369 \pm 0.0029 \rightarrow 0.0279 \pm 0.0009$ on Pleural Effusion). For positive-weighted models, Brier scores are severely elevated and appear to deteriorate over time ($0.1081 \pm 0.0394 \rightarrow 0.1254 \pm 0.0456$ on DenseNet-121; $0.1256 \pm 0.0325 \rightarrow 0.1392 \pm 0.0310$ on ResNet-50), with ECE exceeding 22%. Tracking the excess penalty $\Delta\text{Brier} = \text{Brier}_{\text{weighted}} - \text{Brier}_{\text{unweighted}}$ across sequential strata reveals the operational mechanism: when prevalence declines (Pleural Effusion: 5.27% $\rightarrow$ 3.49%), the excess penalty expands ($+0.0712 \rightarrow +0.0975$); when prevalence rises (Cardiomegaly: 7.82% $\rightarrow$ 10.59%), it attenuates ($+0.0729 \rightarrow +0.0683$); and when prevalence is steady (Pneumonia: 2.66% $\rightarrow$ 2.05%), it remains fixed ($+0.1079 \rightarrow +0.1112$). A static training logit shift produces surplus false-positive probability mass that is penalized more severely when deployment disease prevalence drops, expanding the apparent temporal calibration gap.

3.2. Multi-Calibrator Comparison

Table 3 and Fig. 1 benchmark post-hoc calibrators on Pleural Effusion across 2,000 patient-clustered bootstrap resamples on the 3-seed ensemble.

Table 3: Multi-Calibrator Comparison on Pleural Effusion: Deidentified-Date Ensemble Performance, Calibration Intercept (α), Slope (β), and Paired Bootstrap Contrasts. Complete 12-column factorial evaluation across all intermediate strata and log-loss is provided in Supplementary Table S4.

Architecture	Loss Objective	Calibrator	Brier ($T_1 \rightarrow T_3$)	$T_3 \alpha$	$T_3 \beta$	T_3 ECE	$T_3 \Delta$ Brier vs Raw [95% CI]
DenseNet-121	Weighted BCE	Raw (Uncalibrated)	0.092 $\rightarrow$ 0.108	-2.964	1.048	0.2227	—
DenseNet-121	Weighted BCE	Temperature Scaling (T^*)	0.091 $\rightarrow$ 0.107	-2.962	1.029	0.2190	-0.0011 [-0.0014, -0.0009]
DenseNet-121	Weighted BCE	Analytic Offset ($z - \log w$)	0.036 $\rightarrow$ 0.028	-0.445	0.929	0.0126	-0.0795 [-0.0893, -0.0697]
DenseNet-121	Weighted BCE	Platt Scaling ($a + bz$)	0.036 $\rightarrow$ 0.028	-0.229	1.128	0.0176	-0.0798 [-0.0894, -0.0702]
DenseNet-121	Unweighted BCE	Raw Baseline	0.036 $\rightarrow$ 0.027	+0.147	1.044	0.0040	—
ResNet-50	Weighted BCE	Raw (Uncalibrated)	0.110 $\rightarrow$ 0.123	-2.853	1.950	0.2751	—
ResNet-50	Weighted BCE	Temperature Scaling (T^*)	0.088 $\rightarrow$ 0.101	-2.608	1.184	0.1993	-0.0220 [-0.0241, -0.0198]
ResNet-50	Weighted BCE	Analytic Offset ($z - \log w$)	0.040 $\rightarrow$ 0.029	+1.075	1.405	0.0158	-0.0936 [-0.1022, -0.0844]
ResNet-50	Weighted BCE	Platt Scaling ($a + bz$)	0.038 $\rightarrow$ 0.030	-0.191	1.178	0.0223	-0.0930 [-0.1007, -0.0849]
ResNet-50	Unweighted BCE	Raw Baseline	0.036 $\rightarrow$ 0.027	-0.077	1.095	0.0109	—

For weighted DenseNet-121, temperature scaling optimizes $T^* \approx 1.0$, producing an uncalibrated T_3 Brier score of 0.1067 compared to raw 0.1079 (Δ Brier = -0.0011 [95% CI: $-0.0014, -0.0009$]) while leaving the calibration intercept severely displaced ($\alpha = -2.9615$ vs. raw -2.9638). For ResNet-50, temperature scaling achieves modest variance shrinkage (Δ Brier = -0.0220 [$-0.0241, -0.0198$]), but leaves the underlying intercept discrepancy uncorrected ($\alpha = -2.6080$ vs. raw -2.8526). Temperature scaling is mathematically incapable of adjusting this intercept displacement.

In contrast, subtracting the theoretical logit offset ($z - \ln w$, with $\ln(21.51) = +3.068$) drops T_3 Brier from 0.1079 to 0.0284 (Δ Brier = -0.0795 [95% CI: $-0.0893, -0.0697$], 73.7% error reduction), reducing ECE from 22.3% to 1.26% and shifting the Cox intercept from $\alpha = -2.9638$ to -0.4451 while perfectly preserving the calibration slope ($\beta = 0.9292$). Fitted Offset ($z + a^*$) and Platt Scaling achieve identical performance (T_3 Brier: 0.0289 and 0.0281; Δ Brier = -0.0790 and -0.0798). For ResNet-50, Analytic Offset similarly eliminates 76.3% of probability error ($0.1227 \rightarrow 0.0291$, Δ Brier = -0.0936 [95% CI: $-0.1022, -0.0844$]), confirming that intercept neutralization is architecturally universal.

3.2.1. Empirical Verification of Logit Offset: $\Delta z \approx \ln w$

Evaluating logit differences $\Delta z(x) = \bar{z}_w(x) - \bar{z}_u(x)$ across all 16,604 ensemble evaluations (Fig. 2) confirms that finite neural networks cluster tightly around theoretical $\ln w$: Pleural Effusion ($\ln w = 3.07$) observed mean $\Delta z = 3.29 \pm 0.81$ (median 3.25); Pneumonia ($\ln w = 3.99$) observed 3.83 ± 0.74 (median 3.84); and Cardiomegaly ($\ln w = 2.25$) observed 2.64 ± 1.04 (median 2.60). For Edema ($w = 1091.3, \ln w = 7.00$), acute training scarcity (23




Figure 2: Empirical verification of theoretical logit shift ($z_w - z_u \approx \log w$) across 16,604 held-out deidentified-date image evaluations ($8,302 \text{ images} \times 2 \text{ architectures}$, with Δz computed as the mean difference across the three seed-matched weighted/unweighted model pairs). Kernel density estimates (blue) of observed logit differences between positive-weighted and unweighted models cluster tightly around the theoretical population attractor $\log w$ (dashed red line; observed median in solid green) across Pleural Effusion (a), Pneumonia (b), and Cardiomegaly (c). For Edema (d), extreme training class scarcity ($w > 1000$) induces optimization saturation, bounding empirical shifts near $+3.69$.

276 positive cases) causes optimization saturation, bounding empirical shifts
 277 at $\Delta z = 3.69 \pm 1.47$ (median 3.46) due to L_2 regularization. Equal-mass
 278 reliability diagrams (Fig. 3) demonstrate that the analytic offset realigns
 279 curves tightly with the ideal diagonal across all strata.

280 3.3. Selective Prediction and Workload Triage Simulation

281 Table 4 and Fig. 4 summarize the selective triage simulation on later
 282 held-out stratum T_3 ($N = 2,436$ radiographs, 85 positive Pleural Effusion
 283 cases) under development thresholds. At 100% coverage (full automation),
 284 the raw weighted model incurs 16.0 missed cases (FN_{auto}) with Brier 0.1079.
 285 Calibrated Analytic Offset under matched threshold ($t^* \approx 0.022$) yields




Figure 3: Reliability diagrams for Pleural Effusion (DenseNet-121 3-seed ensemble) across held-out deidentified-date strata (T_1, T_2, T_3). Raw weighted models (red) severely over-predict event probabilities across all bins. Analytic offset calibration (blue) and unweighted models (green) remain tightly calibrated to the ideal diagonal across all deployment windows.

Table 4: Selective Prediction and Reader Referral Triage Simulation on Later Held-Out Stratum T_3 ($N = 2,436$ Radiographs, 85 Positive Pleural Effusion Cases). Evaluated under matched Youden decision thresholds determined on development stratum T_1 across primary clinical operating regimes (100%, 90%, 80%, 70% coverage). Complete 24-row simulation across all coverage levels is detailed in Supplementary Table S5.

Coverage	Referred (Workload)	Strategy	Automated FN (Missed)	Retained Sens.	Full Sens.	Retained AUROC	Retained Brier
100%	0 (0%)	Weighted BCE (Raw)	16.0	81.18%	81.18%	0.8691	0.1079
100%	0 (0%)	Weighted BCE (Analytic Corrected)	16.0	81.18%	81.18%	0.8623	0.0284
100%	0 (0%)	Unweighted BCE (Raw)	12.0	85.88%	85.88%	0.8806	0.0268
90%	243 (10%)	Weighted BCE (Raw)	14.0	82.93%	83.53%	0.8789	0.1062
90%	243 (10%)	Weighted BCE (Analytic Corrected)	15.0	81.01%	82.35%	0.8735	0.0289
90%	243 (10%)	Unweighted BCE (Raw)	12.0	84.62%	85.88%	0.8923	0.0266
80%	487 (20%)	Weighted BCE (Raw)	10.0	87.01%	88.24%	0.8882	0.1012
80%	487 (20%)	Weighted BCE (Analytic Corrected)	15.0	77.61%	82.35%	0.8896	0.0264
80%	487 (20%)	Unweighted BCE (Raw)	12.0	82.61%	85.88%	0.9027	0.0256
70%	730 (30%)	Weighted BCE (Raw)	7.0	89.39%	91.76%	0.8997	0.0956
70%	730 (30%)	Weighted BCE (Analytic Corrected)	13.0	70.45%	84.71%	0.8876	0.0178
70%	730 (30%)	Unweighted BCE (Raw)	11.0	78.00%	87.06%	0.9051	0.0194
70%	730 (30%)	Random Deferral Baseline	11.3	80.99%	86.69%	—	0.1081




Figure 4: Selective prediction evaluation on later held-out stratum T_3 with matched development thresholds (Table 4). Left: Risk-coverage curves demonstrating that selective deferral steadily reduces Brier error across all models as ambiguous cases are referred. Right: Reader referral workload vs. automated false negatives, illustrating the distinct operational behaviors of raw weighted models (halving missed cases to 7.0 under 30% referral) and analytic corrected models (optimizing probability accuracy with Brier 0.0178).

identical sensitivity (81.18%) and 16.0 missed cases, but with an almost four-fold lower Brier score of 0.0284. Unweighted raw models achieve 12.0 missed cases (sensitivity 85.88%, Brier 0.0268).

Under selective deferral at 70% coverage (30% referral workload, 730 radiographs referred): raw weighted models halve automated false negatives from 16.0 to 7.0 (retained sensitivity 89.39%, full-cohort triage sensitivity 91.76%), but leave probabilities uncalibrated (Brier 0.0956); conversely, analytic calibrated models achieve superior probability estimation (Brier 0.0178), but retain 13.0 automated false negatives because threshold-distance normalization ($t^* \approx 0.022$) classifies low-probability false negatives ($\hat{p} \ll t^*$) as confident negatives. This reveals a fundamental operational trade-off:

probability calibration is required for valid risk scoring, whereas raw weighted models with threshold-proximity triage act as a heuristic for recall-oriented screening. Neither system eliminates missed disease entirely.

3.4. Cross-Institutional External Evaluation (MIMIC-CXR-JPG)

The external MIMIC-CXR-JPG cohort comprised 3,403 frontal radiographs from 3,041 studies and 289 patients. Rather than asserting turnkey diagnostic transportability—given moderate discrimination for Pleural Effusion (AUROC 0.688 [95% CI: 0.656, 0.717] DenseNet-121; 0.677 ResNet-50), weak discrimination for Pneumonia (0.583–0.607) and Edema (0.516–0.581), and near-chance discrimination for Cardiomegaly (0.519–0.525)—this experiment functions as an external stress-test of zero-retraining recalibration under domain shift. Negative Brier skill scores (BSS < 0; Table 5) confirmed substantial cross-institutional baseline calibration shift relative to a local null model ($\text{Brier}_{\text{null}} = \pi(1 - \pi)$).

Paired analytic recalibration ($\bar{p}_{\text{corr}} = \frac{1}{3} \sum \sigma(z_s - \ln w)$) demonstrated statistically significant Brier reductions for Cardiomegaly ($\Delta\text{Brier} = -0.0579$ [95% CI: $-0.0790, -0.0349$] DenseNet; -0.0611 [$-0.0818, -0.0415$] ResNet) and Pneumonia ($\Delta\text{Brier} = -0.0095$ [$-0.0170, -0.0016$] DenseNet; -0.0155 [$-0.0254, -0.0059$] ResNet), with point-estimate reductions on Pleural Effusion (-0.0084 and -0.0109). For Edema, subtracting theoretical $\ln w$ worsened Brier score ($+0.0217$), consistent with over-correction when empirical finite-network shifts saturate below theoretical $\ln w$. Diagnostic study-level sensitivity analysis ($N = 3,041$ studies across 289 patients, maximum probability pooling; Supplementary Table S2) demonstrated identical concordance, confirming that recalibration patterns hold across radiographs and studies.

4. Discussion

4.1. Decoupling Loss Distortion from Deidentified-Date Cohort Shift

Our findings clarify the mechanism behind reported “temporal calibration decay” in medical AI [8, 9]. While moderate discriminatory attenuation does occur across subsequent cohort strata (AUROC drops by ≈ 0.06 identically across unweighted and weighted models), the perception of progressive calibration collapse is largely an artifact of positive loss weighting interacting with shifting disease prevalence. Because standard post-hoc evaluation relies almost exclusively on temperature scaling—which modifies logit dispersion without shifting the intercept—evaluators have routinely conflated static

Table 5: Cross-Institutional External Validation on the Official MIMIC-CXR-JPG Frontal Test Cohort ($N = 3,403$ Images from 3,041 Studies Across $N = 289$ Unique Patients). Primary experimental evaluation of paired zero-parameter analytic logit recalibration ($\bar{p}_{\text{raw}} = \frac{1}{3} \sum \sigma(z_s) \rightarrow \bar{p}_{\text{corr}} = \frac{1}{3} \sum \sigma(z_s - \log w)$) on 3-seed ensembles trained with positive-weighted BCE. Evaluated with empirical prevalence (π), local null reference baseline (Brier_{null} = $\pi(1 - \pi)$), patient-clustered 1,000-replicate bootstrap 95% confidence intervals, Brier skill score (BSS = $1 - \text{Brier}/\text{Brier}_{\text{null}}$), and paired contrast ($\Delta\text{Brier} = \text{Brier}_{\text{corr}} - \text{Brier}_{\text{raw}}$). The complete 16-row factorial comparison with unweighted BCE baselines is detailed in Supplementary Table S1, and study-level sensitivity analysis is detailed in Supplementary Table S2.

Architecture	Pathology	Prevalence (π)	Brier _{null}	AUROC [95% CI]	Brier (Raw $\rightarrow$ Corr)	ΔBrier [95% CI]	BSS (Raw $\rightarrow$ Corr)
DenseNet-121	Pleural Effusion	32.18%	0.2182	0.668 [0.636, 0.700]	0.2536 $\rightarrow$ 0.2451	-0.0084 [-0.0401, +0.0240]	-0.162 $\rightarrow$ -0.123
DenseNet-121	Cardiomegaly	26.33%	0.1940	0.525 [0.492, 0.558]	0.2907 $\rightarrow$ 0.2328	-0.0579 [-0.0790, -0.0349]	-0.499 $\rightarrow$ -0.200
DenseNet-121	Pneumonia	10.05%	0.0904	0.598 [0.553, 0.639]	0.1078 $\rightarrow$ 0.0983	-0.0095 [-0.0170, -0.0016]	-0.192 $\rightarrow$ -0.087
DenseNet-121	Edema	21.33%	0.1678	0.581 [0.549, 0.611]	0.1914 $\rightarrow$ 0.2131	+0.0217 [+0.0118, +0.0307]	-0.141 $\rightarrow$ -0.270
ResNet-50	Pleural Effusion	32.18%	0.2182	0.620 [0.588, 0.653]	0.2665 $\rightarrow$ 0.2557	-0.0109 [-0.0417, +0.0253]	-0.221 $\rightarrow$ -0.172
ResNet-50	Cardiomegaly	26.33%	0.1940	0.522 [0.491, 0.552]	0.2822 $\rightarrow$ 0.2211	-0.0611 [-0.0818, -0.0415]	-0.455 $\rightarrow$ -0.140
ResNet-50	Pneumonia	10.05%	0.0904	0.589 [0.545, 0.629]	0.1137 $\rightarrow$ 0.0982	-0.0155 [-0.0254, -0.0059]	-0.257 $\rightarrow$ -0.086
ResNet-50	Edema	21.33%	0.1678	0.570 [0.538, 0.602]	0.1929 $\rightarrow$ 0.2130	+0.0201 [+0.0092, +0.0313]	-0.149 $\rightarrow$ -0.269

objective-loss bias with intrinsic temporal feature drift. When the mathematical logit offset is neutralized ($z - \ln w$), the apparent calibration collapse vanishes, and Brier scores track the unweighted baseline across all held-out strata.

4.2. Relationship to Prior Theoretical and Clinical Literature

Our theoretical framing connects with the work of Caplin, Martin, and Marx [14], who demonstrated that class weighting corrupts calibration and derived an inverse likelihood recovery framework on static single-split tasks. We advance this literature into longitudinal clinical deployment: (i) resolving the temporal calibration decay paradox across sequential deidentified-date cohort strata; (ii) proving the structural failure of temperature scaling ($\alpha \approx -2.96$); (iii) demonstrating finite-network optimization bounds ($w > 1000$ in Edema saturates at $\Delta z \approx 3.69$ vs. theoretical 7.00); and (iv) stress-testing cross-institutional transfer from BRAX to MIMIC-CXR. Furthermore, extending multi-dataset selective CXR investigations [20], we operationalize the clinical triage trade-off: analytic offset ($z - \ln w$) is essential for epidemiological risk stratification and prognosis (restoring Brier parity, consistent with plug-in principles [17]), while raw weighted models support recall-oriented screening triage under referral budgets.

351 4.3. Architectural Universality and Operational Utility

352 The observed phenomena are consistent across architectures: DenseNet-
 353 121 and ResNet-50 exhibit identical discrimination attenuation ($\Delta\text{AUROC} =$
 354 -0.0604 vs. -0.0687), severe raw calibration intercepts ($\alpha \approx -2.96$ and
 355 -2.85), and comparable Brier reductions under Analytic Offset (73.7% and
 356 76.3%, $\Delta\text{Brier} = -0.0795$ and -0.0936). Because Analytic Offset ($z - \ln w$)
 357 operates in constant time ($O(1)$) with zero validation fitting and zero refitting,
 358 it provides an immediate algorithmic solution for resource-constrained edge
 359 devices and PACS pipelines.

360 4.4. Scope, Boundaries, and Methodological Limitations

361 Several methodological limitations define the boundary of these findings:

- 362 1. **Date Anonymization and Temporal Semantics:** In BRAX, dates
 363 are fictitious hashes preserving intra-patient acquisition intervals; cross-
 364 patient calendar synchronization is unverified. Partitions (T_1, T_2, T_3)
 365 reflect sequential deidentified-date strata rather than synchronized cal-
 366 endar years.
- 367 2. **Cross-Institutional Domain Gap and Modest Transportability:**
 368 External validation on MIMIC-CXR-JPG yielded modest discrimina-
 369 tion (Effusion AUROC ≈ 0.69) and near-chance discrimination for
 370 Cardiomegaly (0.52). While analytic offset reduced Brier error for
 371 Cardiomegaly and Pneumonia, it worsened Edema (+0.0217), demon-
 372 strating that zero-retraining offset correction cannot compensate for
 373 severe cross-institutional covariate shift or negative transfer.
- 374 3. **Label Extraction Noise and Retrospective Simulation:** La-
 375 bels were derived from Portuguese reports via NLP rather than multi-
 376 radiologist panel review, and evaluations were retrospective. Real-world
 377 validation requires prospective silent shadow deployment in operational
 378 hospital PACS.
- 379 4. **Architectural Scope and Scarcity Bounds:** We evaluated standard
 380 CNN backbones (DenseNet-121, ResNet-50); Vision Transformers and
 381 vision-language models warrant future investigation. Under extreme
 382 class scarcity (Edema, $w = 1091.3$), empirical logit shifts saturate at
 383 $+3.69$ due to L_2 regularization, defining the empirical boundary of the
 384 asymptotic derivation.

385 5. Conclusion

386 In this controlled study across 40,523 BRAX radiographs, severe probabil-
387 ity miscalibration across held-out strata was predominantly driven by positive
388 loss weighting rather than progressive temporal drift. Because standard
389 temperature scaling cannot shift the logit intercept, it leaves this distortion
390 intact. A constant-time closed-form analytic offset ($z - \ln w$) eliminates
391 73.7% of stratum T_3 Brier error ($0.1079 \rightarrow 0.0284$, $\Delta\text{Brier} = -0.0795$ [95%
392 CI: $-0.0893, -0.0697$]) without model refitting or validation data, restoring
393 calibration parity with unweighted models. Furthermore, selective triage
394 simulations reveal an operational trade-off: raw weighted models provide
395 effective recall-oriented triage by halving missed cases under a 30% refer-
396 ral budget, whereas analytic calibration restores reliable probabilities for
397 decision-theoretic risk stratification. Neither approach eliminates missed
398 disease entirely, underscoring the necessity of transparent, workload-aware
399 deployment policies in clinical deep learning.

400 CRediT Authorship Contribution Statement

401 **Vishesh Panghal:** Conceptualization, Methodology, Software, Validation,
402 Formal Analysis, Investigation, Data Curation, Writing – Original Draft,
403 Writing – Review & Editing, Visualization, Project Administration.

404 Declaration of Competing Interest

405 The author declares that he has no known competing financial interests
406 or personal relationships that could have appeared to influence the work
407 reported in this paper.

408 Data and Code Availability

409 The BRAX dataset (v1.1.0) [2] is available on PhysioNet at <https://physionet.org/content/brax/1.1.0/> under the PhysioNet Credentialed
410 Data Use Agreement. The MIMIC-CXR-JPG database (v2.1.0) [21, 22] is
411 available on PhysioNet at <https://physionet.org/content/mimic-cxr-jpg/2.1.0/> under the PhysioNet Credentialed Data Use Agreement. Access re-
412 quires credentialed PhysioNet researcher status, completion of human subjects
413 research training, and signed data use agreements. Complete end-to-end repli-
414 cation scripts, model evaluation routines, cross-fitting recalibration pipelines,
415
416

selective prediction triage simulations, and tabular evaluation artifacts are
openly available on GitHub at [https://github.com/Vishesh-panghal/
temporal-brax-calibration](https://github.com/Vishesh-panghal/temporal-brax-calibration), and the repository will be permanently as-
signed a persistent DOI via Zenodo upon acceptance of the manuscript. No
credentialed patient-level images, metadata, labels, identifiers, or prediction
rows are redistributed; the repository contains aggregate evaluation artifacts
and a synthetic manifest schema example.

Ethical Approval and Consent to Participate

This study represents a secondary analysis of credentialed-access, de-
identified diagnostic imaging databases accessible under formal Data Use
Agreements. The primary BRAX dataset collection was approved by the
Institutional Review Board (IRB) of Hospital Israelita Albert Einstein (São
Paulo, Brazil; CAAE: 36720520.1.0000.0071) with a formal waiver of informed
consent. Creation and distribution of the MIMIC-CXR-JPG database was
approved by the Institutional Review Boards of Beth Israel Deaconess Medical
Center (Boston, MA) and the Massachusetts Institute of Technology (Protocol
#0403000206) with a waiver of informed consent. The author completed
mandatory human subjects research training (CITI Program certification),
executed approved PhysioNet Data Use Agreements for both databases, and
strictly complied with data governance protocols prohibiting re-identification
or redistribution of protected health information.

Declaration of Generative AI and AI-Assisted Technologies in the Writing Process

During the preparation of this work, the author utilized generative arti-
ficial intelligence platforms (Google Gemini, Anthropic Claude, and OpenAI
Codex) exclusively for LaTeX formatting and syntax verification, analysis code
refactoring and modularization, grammatical refinement, and bibliographic
cross-referencing. Following the use of these tools, the author thoroughly
reviewed, verified, and edited all text, numerical outputs, and code, and takes
full intellectual and scientific responsibility for the integrity, validity, and
conclusions of the published manuscript.

448 Funding

449 This research did not receive any specific grant from funding agencies in
450 the public, commercial, or not-for-profit sectors.

451 References

- 452 [1] P. Rajpurkar, J. Irvin, K. Zhu, B. Yang, H. Mehta, T. Duan, D. Ding,
453 A. Bagul, C. Langlotz, K. Shpanskaya, et al., CheXNet: Radiologist-level
454 pneumonia detection on chest X-rays with deep learning, arXiv preprint
455 arXiv:1711.05225 (2017).
- 456 [2] E. P. Reis, J. P. Q. de Paiva, M. C. B. da Silva, G. A. S. Ribeiro,
457 V. F. Paiva, L. Bulgarelli, H. M. H. Lee, P. V. Santos, V. M. Brito,
458 L. T. W. Amaral, G. L. Beraldo, J. N. Haidar Filho, G. B. S. Teles,
459 G. Szarf, T. Pollard, A. E. W. Johnson, L. A. Celi, E. Amaro, BRAX, a
460 Brazilian labeled chest X-ray dataset, Scientific Data 9 (1) (2022) 487.
461 [doi:10.1038/s41597-022-01608-8](https://doi.org/10.1038/s41597-022-01608-8).
- 462 [3] S. G. Finlayson, A. Subbaswamy, K. Singh, J. Bowers, A. Kupke, J. Zit-
463 train, I. S. Kohane, S. Saria, The clinician and dataset shift in artificial
464 intelligence, New England Journal of Medicine 385 (3) (2021) 283–286.
465 [doi:10.1056/NEJMc2104626](https://doi.org/10.1056/NEJMc2104626).
- 466 [4] A. Kore, E. A. Babil, V. Subasri, M. Abdalla, B. Fine, E. Dolatabadi,
467 M. Abdalla, Empirical data drift detection experiments on real-world
468 medical imaging data, Nature Communications 15 (1) (2024) 1887. [doi:](https://doi.org/10.1038/s41467-024-46142-w)
469 [10.1038/s41467-024-46142-w](https://doi.org/10.1038/s41467-024-46142-w).
- 470 [5] M. Shenouda, I. Flerlage, A. Kaveti, M. L. Giger, S. G. Armato, III,
471 Assessment of a deep learning model for COVID-19 classification on
472 chest radiographs: a comparison across image acquisition techniques
473 and clinical factors, Journal of Medical Imaging 10 (6) (2023) 064504.
474 [doi:10.1117/1.jmi.10.6.064504](https://doi.org/10.1117/1.jmi.10.6.064504).
- 475 [6] B. Van Calster, D. J. McLernon, M. van Smeden, L. Wynants, E. W.
476 Steyerberg, Calibration: the Achilles heel of predictive analytics, BMC
477 Medicine 17 (1) (2019) 230. [doi:10.1186/s12916-019-1466-7](https://doi.org/10.1186/s12916-019-1466-7).

- 478 [7] E. W. Steyerberg, A. J. Vickers, N. R. Cook, T. Gerds, M. Gonen,
479 N. Obuchowski, M. J. Pencina, M. W. Kattan, Assessing the performance
480 of prediction models: a framework for traditional and novel measures,
481 *Epidemiology* 21 (1) (2010) 128–138.
- 482 [8] L. L. Guo, S. R. Pfohl, J. Fries, A. E. W. Johnson, J. Posada,
483 C. Aftandilian, N. Shah, L. Sung, Evaluation of domain general-
484 ization and adaptation on improving model robustness to temporal
485 dataset shift in clinical medicine, *Scientific Reports* 12 (1) (2022) 2726.
486 [doi:10.1038/s41598-022-06484-1](https://doi.org/10.1038/s41598-022-06484-1).
- 487 [9] J. S. Hinson, E. Y. Klein, A. Smith, et al., Multisite implementation
488 of a workflow-integrated machine learning system to optimize COVID-
489 19 hospital admission decisions, *npj Digital Medicine* 5 (1) (2022) 94.
490 [doi:10.1038/s41746-022-00646-1](https://doi.org/10.1038/s41746-022-00646-1).
- 491 [10] M. Buda, A. Maki, M. A. Mazurowski, A systematic study of the class
492 imbalance problem in convolutional neural networks, *Neural Networks*
493 106 (2018) 249–259.
- 494 [11] A. Alexandari, A. Kundaje, A. Shrikumar, Maximum likelihood with
495 bias-corrected calibration is hard-to-beat at label shift adaptation, *Inter-
496 national Conference on Machine Learning (ICML)* (2020) 222–232.
- 497 [12] S. E. Davis, R. A. Greevy, C. Fonnesbeck, T. A. Lasko, C. G. Walsh,
498 M. E. Matheny, Detection of calibration drift in clinical prediction models
499 to inform model updating, *Journal of Biomedical Informatics* 112 (2020)
500 103611.
- 501 [13] G. King, L. Zeng, Logistic regression in rare events data, *Political Analysis*
502 9 (2) (2001) 137–163.
- 503 [14] A. Caplin, D. Martin, P. Marx, Calibrating for class weights by modeling
504 machine learning, *arXiv preprint arXiv:2205.04613* (2022).
- 505 [15] C. Guo, G. Pleiss, Y. Sun, K. Q. Weinberger, On calibration of modern
506 neural networks, *International Conference on Machine Learning (ICML)*
507 (2017) 1321–1330.

- 508 [16] J. Platt, Probabilistic outputs for support vector machines and com-
 509 parisons to regularized likelihood methods, *Advances in Large Margin*
 510 *Classifiers* 10 (3) (1999) 61–74.
- 511 [17] M. Piccininni, M. Wechsung, B. Van Calster, J. L. Rohmann, S. Konig-
 512 orski, M. van Smeden, Understanding random resampling techniques
 513 for class imbalance correction and their consequences on calibration and
 514 discrimination of clinical risk prediction models, *Journal of Biomedical*
 515 *Informatics* 155 (2024) 104666. doi:10.1016/j.jbi.2024.104666.
- 516 [18] Y. Geifman, R. El-Yaniv, Selective classification for deep neural networks,
 517 *Advances in Neural Information Processing Systems (NeurIPS)* 30 (2017)
 518 4878–4887.
- 519 [19] B. Kompa, J. Snoek, A. L. Beam, Second opinion needed: communicating
 520 uncertainty in medical machine learning, *npj Digital Medicine* 4 (1) (2021)
 521 4. doi:10.1038/s41746-020-00367-3.
- 522 [20] Y. Aperstein, T. Amit, G. Alon, V. Tal, S. D. Ravit, A. Apartsin,
 523 Knowing when not to decide: multi-pathology chest X-ray classification
 524 with rejection mechanisms, *Scientific Reports* 16 (2026) 66294. doi:
 525 10.1038/s41598-026-66294-7.
- 526 [21] A. E. W. Johnson, T. J. Pollard, S. J. Berkowitz, N. R. Greenbaum, M. P.
 527 Lungren, C.-y. Deng, R. G. Mark, S. Horng, MIMIC-CXR, a de-identified
 528 publicly available database of chest radiographs with free-text reports,
 529 *Scientific Data* 6 (1) (2019) 317. doi:10.1038/s41597-019-0322-0.
- 530 [22] A. Johnson, M. Lungren, Y. Peng, Z. Lu, R. Mark, S. Berkowitz, S. Horng,
 531 MIMIC-CXR-JPG - chest radiographs with structured labels, <https://physionet.org/content/mimic-cxr-jpg/2.1.0/> (2024). doi:10.
 532 13026/jsn5-t979.
- 534 [23] G. Huang, Z. Liu, L. Van Der Maaten, K. Q. Weinberger, Densely
 535 connected convolutional networks, *Proceedings of the IEEE Conference*
 536 *on Computer Vision and Pattern Recognition (CVPR)* (2017) 4700–4708.
- 537 [24] K. He, X. Zhang, S. Ren, J. Sun, Deep residual learning for image
 538 recognition, *Proceedings of the IEEE Conference on Computer Vision*
 539 *and Pattern Recognition (CVPR)* (2016) 770–778.